


Transport layer: roadmap


- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer: Go-Back-N, SR
 - reliable data transfer: TCP
 - TCP connection management
 - TCP flow control
 - TCP congestion control

Recap: TCP flow control


- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
 - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
- **sender** limits amount of unACKed (“in-flight”) data to received **rwnd**
$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{rwnd}$$
- guarantees receive buffer will not overflow



Flow vs congestion control



TCP congestion control



TCP sending behavior:


- *roughly*: send $cwnd$ bytes, wait RTT for ACKS, then send more bytes

$$\text{TCP rate} \approx \frac{cwnd}{RTT} \text{ bytes/sec}$$

- TCP **sender** limits transmission: $\text{LastByteSent} - \text{LastByteAcked} \leq cwnd$
- $cwnd$ is dynamically adjusted in response to observed network congestion
 - initially, $cwnd = 1$ MSS (maximal segment size)

Slow start

- when connection begins, increase rate exponentially until first loss event
 - on each new ack
 - $cwnd += 1 \text{ MSS}$
 - effectively, doubling $cwnd$ every RTT
 - “start small, but grow really fast”
 - Q: why?



Congestion avoidance

- Congestion avoidance
 - Slow start threshold: $ssthresh$
 - when $cwnd > ssthresh$
 - on each new ack
 - $cwnd += \frac{MSS^2}{cwnd}$
 - effectively, $cwnd += 1 \text{ MSS}$ every RTT
 - linear increment
- Q: why increase $cwnd$ to “avoid” congestion?


TCP: from slow start to congestion avoidance

Q: when should the exponential increase switch to linear?

A: when **cwnd** gets to 1/2 of its value before timeout.

Implementation:

- variable **ssthresh**
- on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event



https://media.pearsoncmg.com/ph/esm/ecs_kurose_compnetwork_8/cw/content/interactiveanimations/tcp-congestion/index.html

Network congestion

- cwnd is always increased in slow-start and congestion avoidance
 - network congestion is inevitable
- Network congestion indicator
 - TCP treats packet loss as network congestion
- Packet loss indicators
 - acknowledgment timeout
 - 3 duplicate acknowledgments


Q: packet loss == network congestion?

Timeout retransmit

- Congestion control
 - $\text{ssthresh} = \text{cwnd} / 2$
 - $\text{cwnd} = 1 \text{ MSS}$
 - followed by slow-start
- Error control
 - retransmit packet
 - restart timer

Congestion window

- Slow-start
- Congestion avoidance
- Timeout retransmission
 - TCP timeout is quite conservative
 - pay attention to how ssthresh is adjusted!




Fast retransmit

- Duplicate acknowledgment
 - example
 - rcv: [0, 499], [500, 999], [1500, 1999], [2000, 2499], [2500, 2999]
 - ack: 500, 1000, 1000, 1000, 1000 (3rd dupack)
- Congestion control (fast retransmit)
 - on 3rd dupack: $\text{ssthresh} = \text{cwnd}/2$; $\text{cwnd}=1 \text{ MSS}$
 - followed by slow start
- Error control
 - retransmit: [1000, 1499]

Q: dupack == loss?

Fast retransmit: cwnd




Q: why “fast” retransmit?

TCP Reno: Fast recovery

- Duplicate acknowledgment
 - example
 - rcv: [0, 499], [500, 999], [1500, 1999], [2000, 2499], [2500, 2999]
 - ack: 500, 1000, 1000, 1000, 1000 (3rd dupack)
- Congestion control (fast recovery)
 - on 3rd dupack: $ssthresh=cwnd/2$, $cwnd= ssthresh + 3$
 - followed by congestion avoidance
- Error control
 - retransmit: [1000,1499]


Q: fast transmit vs recovery?

Fast recovery: cwnd



Q: why fast recovery?

Summary: TCP Reno congestion control



TCP congestion control: AIMD


- *approach:* senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event

Additive Increase

increase sending rate by 1 maximum segment size every RTT until loss detected

Multiplicative Decrease

cut sending rate in half at each loss event



AIMD sawtooth behavior: *probing* for bandwidth

TCP AIMD: more

Multiplicative decrease detail: sending rate is

- Cut in half on loss detected by triple duplicate ACK (TCP Reno)
- Cut to 1 MSS (maximum segment size) when loss detected by timeout (TCP Tahoe)

Why AIMD?

- AIMD – a distributed, asynchronous algorithm – has been shown to:
 - optimize congested flow rates network wide!
 - have desirable stability properties

Questions: TCP Reno

- Num timeouts?
- Num 3-duplicate acks?
- Num of congestion avoidance phases?
- Number of slow start phases?
- Verify cwnd values
- sshthresh values?

